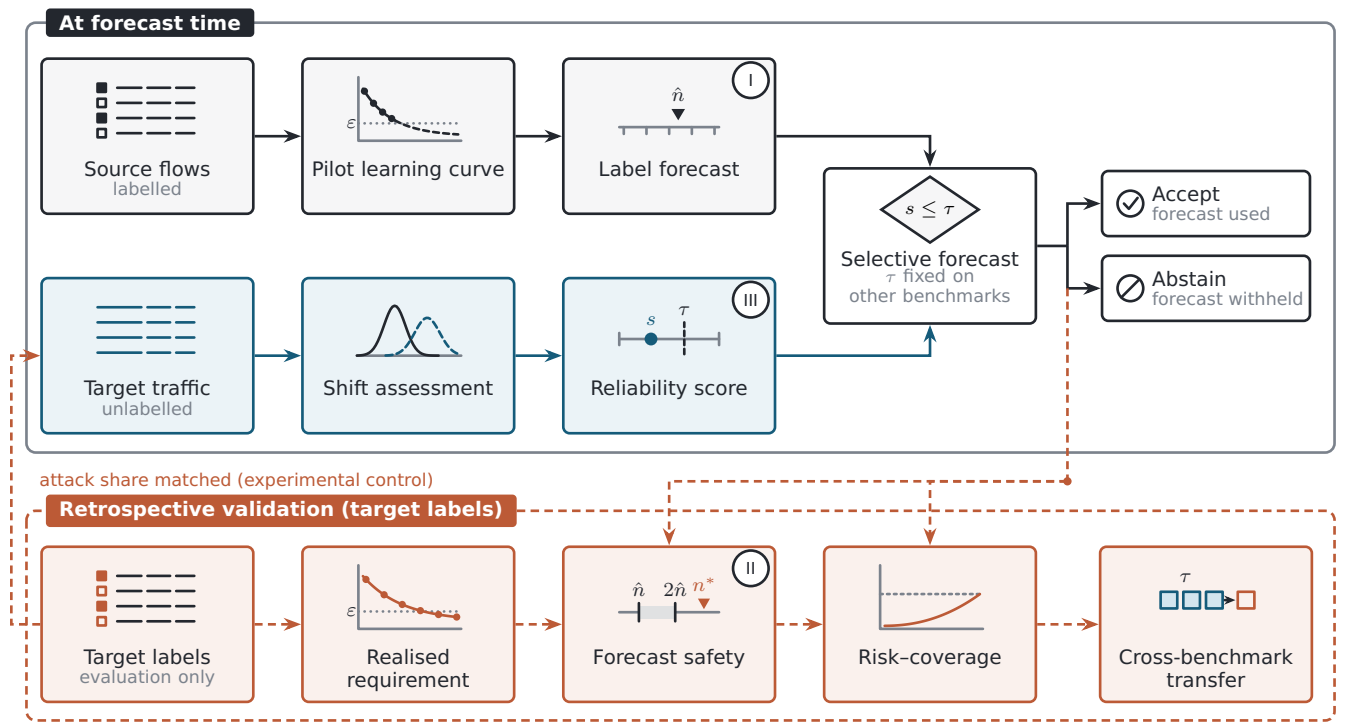
At forecast time
Source flows
labelled
Pilot learning curve
ε
$\hat{n}$
Label forecast
I
$s \le \tau$
Selective forecast
τ fixed on
other benchmarks
Accept
forecast used
Abstain
forecast withheld
Target traffic
unlabelled
Shift assessment
s
τ
Reliability score
III
attack share matched (experimental control)
Retrospective validation (target labels)
Target labels
evaluation only
Realised
requirement
$\hat{n}$
$2\hat{n}$
n^*
Forecast safety
II
Risk–coverage
τ
Cross-benchmark
transfer